


## Topic 2: Protecting the OS



Reading: Sections 2.1-2.4

Next reading: same

- Think of the OS as a library that contains routines<sup>1</sup> that everyone needs. Consider a routine that reads a file from the disk. How does the OS ensure that the file is protected properly, e.g. only the owner can read it?
  - One possibility is to check permissions at the beginning of the routine, and return if the read is not allowed.
  - What prevents someone from writing assembly code that jumps into the routine after the protection is checked?
  - Procedure calls offer no protection.
- The OS needs hardware support to protect itself from applications.
  - Use user/supervisor (*kernel*) mode of processor.
  - Certain machine instructions can only be executed while in kernel mode (*privileged instructions*).
  - Some of these instructions allow memory to be made off-limits to user mode. The OS puts itself this memory.


---

<sup>1</sup> I will use the terms procedure, function, routine, and subroutine interchangeably.


- User programs run in user mode; OS runs in kernel mode.  
For this reason the OS is often called the kernel.
- Processor switches from user mode to kernel mode when an **interrupt or trap** occurs.
  - User program cannot switch to kernel mode directly!
- OS uses privileged instruction to switch back to user mode once the interrupt or trap has been handled.
- Processor powers-up in kernel mode.
- *Interrupt*: asynchronous event that causes the processor to stop what it is doing and begin executing at a specified address. Has nothing to do with the instruction being executed, e.g. disk interrupt. *<Keyboard Interrupt>*
  - CPU switches to kernel mode if it is in user mode.
  - Routine located at the address is called an *interrupt handler*.
  - *Interrupt vector* is a table of interrupt handlers, one for each type of interrupt.
  - Processor switches to a separate *interrupt stack* or *kernel stack*.


- Reliability

- Security



- As part of invoking an interrupt handler the processor *masks* interrupts of the same type. OS can change an interrupt's mask.
- Interrupts sometimes have priorities, depending on the CPU.
  - Higher priority interrupts can interrupt lower-priority *open file* handlers.
- *Trap*: synchronous event associated with the execution of the current instruction, e.g. divide-by-zero, overflow.
  - A *system call* is a special kind of trap. It is not an error, but a request by the user program for the OS to do something. Typically processors have a trap instruction that is used to implement system calls.
- OS is interrupt-driven. *The only way to enter it is via an interrupt or trap.*
- User programs can only access OS through system calls (traps), allowing the OS to protect itself. *the only way to switch from user to sys (kernel)*
- Interrupts make the OS *non-deterministic*. It is difficult to deal with so many things going on at one time.
- OS software structure
  - Monolithic (e.g. Linux)





- OS is a protected library. User programs invoke system calls via trap. OS services implemented as a bunch of procedures.
- Some OS-related programs never leave the kernel. More on this later.
- Can be messy. Requires good software engineering practices.
- More on other OS architectures later.