

Topic 5: Synchronization: The Too-Much-Milk Problem

Reading: Section 5.1

Next reading: 5.1-5.5

- The “too much milk” problem.

Roommate A	Roommate B
3:00	
Look in fridge. Out of milk.	
3:05	
Leave for store.	
3:10	
Arrive at store.	Look in fridge. Out of milk.
3:15	
Leave store.	Leave for store.
3:20	
Arrive home, put milk away.	Arrive at store.
3:25	
	Leave store.
3:30	
	Arrive home. OH, NO!

- What does correct mean? Somebody gets milk, but we don’t get too much milk.
- *The most important thing in synchronization is to figure out **precisely** what you want to achieve.*
 - **Safety**: the program never enters a bad state.
 - **Liveness**: the program eventually enters a good state.
Opposite of liveness is *starvation*.
 - E.g., what are the constraints on a traffic light?

- Roommates as a program:

```
1. if (Milk == 0) {  
2.     Milk++;  
3. }
```

- *Synchronization*: using atomic operations to ensure correct cooperation (safety) between processes (eliminate *race conditions*).
 - E.g. buyer and drinker of milk; cars going through an intersection.
 - Does not rely on scheduling priorities.
 - Also known as *inter-process communication (IPC)*.
 - *Atomic operation* -- an operation that is indivisible by other operations.
 - Alternatively, an operation whose effects can be observed either before or after the operation, but not during.
- *Mutual exclusion*: Mechanisms that ensure that only one person or process is doing certain things at one time (others are excluded). E.g. only one person goes shopping at a time.
- *Critical section*: Code in which only one process may be executing simultaneously. E.g. shopping, intersection.
 - Note: critical sections are independent from functions -- a single function may have several critical sections, or a critical

section may span several functions. It is often desirable to have them map one-to-one to make the program more understandable, but it is not required.

- *Mutual exclusion* is the mechanism, a *critical section* is the result.
- There are many ways to achieve mutual exclusion.
 - In kernel mode, you can disable interrupts (uniprocessor only) and not invoke the dispatcher.
 - In user mode, most solutions involve some sort of *locking* mechanism: prevent a process from doing something. If a process tries to acquire a lock that is already held it must either wait, or do something else that doesn't require the lock.
- Three elements of locking:
 1. Lock before using. (leave note)
 2. Unlock when done. (remove note)
 3. Wait (or skip) if locked. (don't shop if note)
- First attempt at computerized milk buying:

Processes A & B

```
1. if (Note == FALSE) {  
2.     if (Milk == 0) {  
3.         Note = TRUE;  
4.         Milk++;  
5.         Note = FALSE;  
6.     }  
7. }
```

- This doesn't always work: A1 A2 B1 B2 B3 A3 B4 A4...
- This solution works for people because lines 1 - 3 are performed atomically: you'll see the other person at the refrigerator and make arrangements. Typically, computers can't both test (look for note) and set (leave note) at the same time.
 - Looking for the note requires a memory read, and leaving a note requires a memory write, which are separate instructions.
 - In this case, we haven't eliminated the problem; we've just moved it and made it a little less likely (i.e. a little more insidious). This is typical of first attempts at solutions to synchronization problems.
 - What happens if we leave the note before looking for a note: does this make everything work? No, no-one will buy milk.
 - What if we check for milk after leaving the note? Still incorrect (A1 B1 A2 B2). In this case, we haven't eliminated