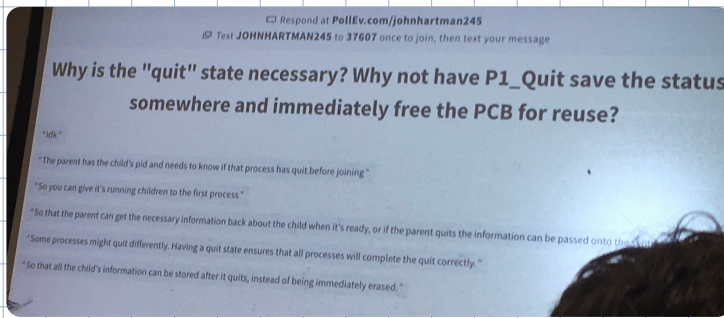


Topic 6: Synchronization via Semaphores

Reading: 5.1-5.5

Next reading: 5.8

- The too-much-milk solution is much too complicated and too limiting. Turning off interrupts isn't possible at user-level. The mutual exclusion mechanism is too simple-minded: it uses only atomic reads and writes. This is sufficient, but unpleasant. It would be unbearable to extend that mechanism to many processes. Let's look at more powerful, higher-level mechanisms.
- *Required properties* of a mutual exclusion mechanism:
 - Must allow only one process into a critical section at a time.
 - Must handle an arbitrary number of processes.
 - If several processes arrive at once, must allow one and only one process to proceed.
 - If one or more processes are waiting to enter the critical section one must do so when the process leaves the critical section (liveness).
- *Desirable properties* of a mutual exclusion mechanism:
 - **Fair**: if several processes waiting, each should get in eventually.



① freeing PCB (status) still in use

② PCB 1 → Quit → reuse

PCB 11 can be reused before completely cleared
PCB11 function might be confused / unexpected result.

- **Efficient**: should not use up substantial amounts of resources when waiting. E.g. no busy waiting.
- **Simple**: should be easy to use (e.g. just bracket the critical sections, lock abstraction).
- Requirements of processes using the mechanism:
 - Always lock before manipulating shared data.
 - Always unlock after manipulating shared data.
- Desirable properties of processes using the mechanism:
 - Do not lock again if already locked.
 - Some mechanisms will cause the process to *deadlock* if this happens, others allow it.
 - Do not unlock if not locked by you
 - There are a few exceptions to this.
 - Do not spend lots of time in critical section.
- *Semaphore*: A synchronization variable that takes on positive integer values. Invented by E. Dijkstra in the mid 60's.
 - **P(semaphore)**: an atomic operation that **waits** for semaphore to become positive, then decrements it by 1 ('Proberen' in Dutch).
 - **V(semaphore)**: an atomic operation that increments semaphore by 1. It does **not** wait. ('Verhoog' in Dutch)

- Semaphores are simple and elegant and allow the solution of many interesting problems. They do a lot more than just mutual exclusion.
- Book doesn't like them, prefers higher-level constructs based on locks and condition variables.
- Too much milk problem with semaphores:

Process A & B

Wait until positive. -1 then

P (OKToBuyMilk) ;	P (OKToBuyMilk) ;
if (Milk == 0) {	if (Milk == 0) {
Milk++;	Milk++;
}	}
V (OKToBuyMilk) ;	V (OKToBuyMilk) ;

+

- Note: OKToBuyMilk must initially be set to 1. What happens if it isn't?
- Why can there can never be more than one process buying milk at once?
- Semaphores aren't provided by hardware, (I'll describe implementation later) but they have several attractive properties:
 - Machine independent.
 - Simple.
 - Work with many processes.